Satwik Medipalli

AI Engineer — LLM Applications, RAG & Multi-Agent Systems

Richardson, TX | 945-544-2418 | sathwik.medipalli@gmail.com | linkedin.com/in/medipalli-satwik | github.com/satwikmed | satwikmedipalli.dev

Summary

AI engineer with four live LLM products built end to end while completing an M.S. at UT Dallas (2026): a multi-agent due-diligence platform, an LLM-judge evaluation system with statistical calibration, and two ML-driven sports analytics apps, plus three internships building NLP pipelines and cloud ML deployments. Every project below is live at a public URL.

Deployed Projects (built during M.S., 2024 – 2026)

Diligence AI: Multi-Agent Due Diligence Platform | *LangGraph, CrewAI, Pinecone* | diligence-ai.vercel.app — Spring 2026

- Built an autonomous due-diligence platform where six specialized agents (LangGraph, CrewAI, OpenAI Agents SDK, Pydantic AI) read SEC 10-K filings, extract financials, classify risk factors, and assemble an analyst-style report in minutes
- Implemented Pinecone vector search with a RAG follow-up interface so users can ask natural-language questions over analyzed filings and receive answers with cited sources; answer faithfulness is scored with RAGAS before every release ships

VoiceIQ: Evaluation Platform for Voice AI Agents | *Python, LLM-as-judge, Streamlit* | voice-agent.vercel.app — Fall 2025

- Created a testing platform that simulates customer conversations across 10 scenario types using a two-LLM loop, scoring agent performance on five dimensions with weighted scoring and self-consistency tracking across repeated runs of each scenario
- Validated the LLM judge itself with Bland-Altman agreement analysis, Pearson/Spearman correlation, and precision/recall tracking against reference scores, backed by 70 unit tests and a pre-seeded database so the platform deploys instantly

NFL Draft Intelligence | *Python, XGBoost, SHAP, SQL* | nfl-draft-intelligence.vercel.app — Spring 2025

- Trained position-specific XGBoost classifiers on two decades of NFL Combine data to evaluate draft prospects, with SHAP values explaining which measurables drive each prediction and a cosine-similarity engine matching prospects to historical comps

NFL Edge: Front-Office & Coaching Analytics | *Python, SQL, Next.js* | nfl-edge.vercel.app — Fall 2024

- Designed three analytics engines grading play-calling predictability, roster efficiency (EPA per salary-cap dollar), and in-game decision quality across all 32 NFL teams, deployed as a production web app with interactive dashboards on Vercel

Experience

AI Research Intern — Feb 2024 – Apr 2024

Boltzmann Labs

- Wrote a format-aware Python preprocessing engine that auto-detects three types of incoming biological data across ∼500K records at 85%+ accuracy and routes each through the correct pipeline, replacing a fully manual data-preparation workflow
- Containerized the ML backend with Docker and deployed it on Kubernetes and AWS S3 behind REST APIs, turning a local prototype into a scalable cloud service that the frontend team could integrate against without touching the model code
- Analyzed competitor biotech analytics platforms and translated findings into product requirements that shipped in BoltBio

Data & AI Intern — Sept 2023 – Dec 2023

Bostonlogix

- Built an NLP extraction pipeline with OpenAI APIs on Microsoft Azure that pulls structured fields from ∼100K unstructured text records, fully automating a manual review process and saving hours of analyst effort every single week
- Migrated ∼1M records from MongoDB to SQL Server with Python to create a queryable analytics store, and audited and restored roughly 1,000 failing API endpoints with Postman that were silently breaking downstream client integrations
- Designed Bold BI dashboards for Hungry Jack's that leadership used in shareholder reviews of store performance

Data Analytics Intern — Apr 2023 – Jun 2023

Ideabytes

- Built a real-time ingestion pipeline that streamed ∼500K JSON sensor datapoints through APIs into Power BI, and presented the live dashboard directly to the CEO to inform decisions on sensor deployment and resource allocation
- Developed linear regression and ARIMA models forecasting sensor performance at 92–95% one-month accuracy and 80–90% at twelve months, letting the team schedule proactive maintenance before hardware failures actually occurred

Education

University of Texas at Dallas — Richardson, TX

M.S. Business Analytics — *May 2026*

GITAM University — Hyderabad, India

B.Tech Computer Science with Data Science — *May 2024*

Technical Skills

LLM & Agents: LangChain, LangGraph, CrewAI, OpenAI Agents SDK, Pydantic AI, RAG, Pinecone, RAGAS, LLM evaluation
ML: XGBoost, SHAP, Scikit-learn, BERT/Transformers, forecasting (ARIMA, linear regression), model explainability
Engineering: Python, SQL, JavaScript, FastAPI, Flask, Docker, Kubernetes, AWS S3, Microsoft Azure, Git, Postman
Data & Visualization: Pandas, NumPy, PostgreSQL, SQL Server, MongoDB, Power BI, Tableau, Bold BI, Streamlit, Plotly