

# Satwik Medipalli

AI Engineer / LLM Engineer | Multiagent Systems, RAG & Production Software  
Livingston, NJ | 945-544-2418 | sathwik.medipalli@gmail.com | LinkedIn | GitHub | Portfolio

## SUMMARY

UT Dallas M.S. Business Analytics (May 2026) builder with three shipped internships and four live AI products: multiagent due diligence over SEC filings, calibrated LLM judge evaluation, a grounded S&P 500 research terminal, and fleet reliability ML with a claim verifying AI analyst. Every project is public and production deployed.

## DEPLOYED PROJECTS

### Diligence AI: Multiagent Due Diligence Platform

Spring 2026

*LangGraph, CrewAI, Pinecone, RAG | Live Demo*

- Built a six agent system that reads SEC 10K filings, extracts financials, classifies risk factors, and assembles citation backed analyst reports in minutes, with Pinecone RAG Q&A and RAGAS faithfulness gating every release

### VoiceIQ: Evaluation Platform for Voice AI Agents

Fall 2025

*Python, LLM as judge, Streamlit | Live Demo*

- Shipped a testing harness that simulates conversations across scenario types, scores agents on five dimensions with self consistency checks, and validates the LLM judge with Bland Altman analysis plus precision/recall against reference scores

### Lumen: S&P 500 Research Terminal with Grounded AI

Summer 2026

*React, TypeScript, FastAPI, Python | Live Demo*

- Built a full stack market terminal on live Yahoo quotes, SEC XBRL, and FRED data with a risk lab, portfolio backtests, and an AI layer constrained to explain only backend computed evidence with no invented numbers or causes

### SiliconSentinel: Fleet Reliability with AI Analyst

Fall 2025

*Python, XGBoost, Pydantic AI | Live Demo*

- Delivered a reliability platform scoring 30 day device failure risk at 83% recall and 80% precision with survival models, SHAP deep dives, an ops RMA queue, and a Pydantic AI analyst that verifies model claims before operators act

## EXPERIENCE

### AI Research Intern

Jan 2024 - Apr 2024

*Boltzmann Labs*

- Built a format aware Python preprocessing engine that auto detects three biological data types across about 500K records at 85%+ accuracy and routes each file through the correct pipeline, replacing a manual prep workflow for the research team
- Turned a local ML prototype into a production cloud service by containerizing the backend with Docker, deploying on Kubernetes with AWS S3 storage, and exposing REST APIs consumed daily by the product and frontend teams

### Data Analytics Intern

Aug 2023 - Dec 2023

*Bostonlogix*

- Shipped an Azure OpenAI NLP pipeline that extracts structured fields from about 100K unstructured text records and fully automates a previously manual review process for operations
- Migrated about 1M records into a SQL Server analytics store and audited and restored roughly 1000 failing API endpoints using Postman test suites across the platform

### BI Developer Intern

Apr 2023 - Jul 2023

*Ideabytes*

- Built a real time ingestion pipeline streaming about 500K JSON sensor datapoints through APIs into Power BI, then presented the live operations dashboard directly to the CEO for review
- Developed regression and ARIMA models forecasting sensor performance at 92-95% one month accuracy and 80 to 90% at twelve months to schedule proactive maintenance before failures

## EDUCATION

### University of Texas at Dallas

Richardson, TX

*M.S. Business Analytics*

*May 2026*

### GITAM University

Hyderabad, India

*B.Tech Computer Science with Data Science*

*May 2024*

## TECHNICAL SKILLS

**LLM & Agents:** LangChain, LangGraph, CrewAI, OpenAI APIs, Pydantic AI, RAG, Pinecone, RAGAS, LLM as judge, prompt workflows, citation grounding

**Software Engineering:** Python, JavaScript/TypeScript, React, FastAPI, Flask, REST APIs, Docker, Kubernetes, Git, Postman

**ML & Evaluation:** Scikit learn, XGBoost, SHAP, survival analysis, model validation, Bland Altman, precision/recall

**Data & Cloud:** SQL, PostgreSQL, SQL Server, MongoDB, AWS S3, Azure, Pandas, NumPy, Streamlit