

Satwik Medipalli

AI / LLM Engineer | RAG, Agents & Eval

Livingston, NJ | 945-544-2418 | sathwik.medipalli@gmail.com | [LinkedIn](#) | [GitHub](#) | [Portfolio](#)

EXPERIENCE

AI Researcher (UTD Client Project)

Jan 2026 – May 2026

DFW Technology

Richardson, TX

- Completed a semester-long UTD educational engagement on-site at DFW Technology, authoring *Public vs Private AI: Scalable Business Agents for SMEs* on Private AI-first architecture, governance, and compliance.
- Authored *When to Use Small Language Models: A Strategic Playbook*, including a Unit of Work pilot measuring >90% inference-cost reduction vs LLM-only at matched task quality.

AI Research Intern

Jan 2024 – Apr 2024

Boltzmann Labs

Hyderabad, India

- Built a multi-omic analysis pipeline over millions of records using clustering, differential expression analysis, and GSEA to prioritize cancer-linked cell populations for downstream validation.
- Containerized and deployed the pipeline with Docker, FastAPI, and Kubernetes, cutting researcher setup from hours to minutes.

Data Analytics Intern

Aug 2023 – Dec 2023

Bostonlogix

Hyderabad, India

- Benchmarked Azure OpenAI against classical NLP for recipe extraction on 100K Hungry Jack's records using precision and recall; the team adopted the LLM path in production, replacing manual ingredient review.
- Diagnosed and fixed 100+ failing API test cases in Postman, added auth and daily scheduled runs, documented APIs in Swagger, and presented the regression suite for release testing.

BI Developer Intern

Apr 2023 – Jul 2023

Ideabytes

Hyderabad, India

- Forecasted one-year performance across 100+ IoT sensors with ARIMA in Python, shifting maintenance from reactive to scheduled.
- Built live Power BI dashboards from IoT JSON APIs for the CEO and the company website, replacing manual weekly reporting.

PROJECTS

DiligenceAI | [Live Demo](#)

Spring 2026

LangGraph, CrewAI, Pinecone, RAGAS, FastAPI, Next.js

- Orchestrated a six-agent A2A pipeline that turns a 10-K into a citation-backed ER report with QoQ filing delta and earnings-vs-filing contradiction checks, compressing first-pass diligence into one session.
- Public History includes complete CRM, MSFT, and AAPL runs; follow-up Q&A is RAG-backed and scored with RAGAS (faithfulness, relevancy, context precision).

VoiceIQ | [Live Demo](#)

Fall 2025

Python, LLM-as-Judge, Next.js, EVA-Bench

- Built a pre-launch eval portal that runs 9 caller scenarios, scores five dimensions with turn-level evidence, and holds deploy below a 90% pass rate at score ≥ 70 .
- Calibrated the judge against a human-graded set (MAE 3.7 pts, Bland-Altman bias +2.8) so gate decisions rest on measured agreement before production.

Lumen | [Live Demo](#)

Summer 2025

React, TypeScript, FastAPI, D3, Yahoo, SEC EDGAR, FRED

- Shipped an S&P 500 research terminal over 503 names with live quotes, SEC XBRL (12 ratios), FRED macro, and AI that only narrates backend-computed evidence so filings are quoted, never recalled.
- Mapped 955 relationship edges across the full index; a Sources page refuses invented numbers when a feed is missing or stale.

EDUCATION

University of Texas at Dallas

Aug 2024 – May 2026

M.S. Business Analytics

Richardson, TX

GITAM University

Aug 2020 – May 2024

B.Tech Computer Science with Data Science

Hyderabad, India

TECHNICAL SKILLS

LLM & Agents: LangGraph, LangChain, CrewAI, Pydantic AI, RAG, OpenAI, Azure OpenAI, Anthropic, LLM-as-Judge

Eval & Retrieval: Pinecone, RAGAS, ChromaDB

Languages & Software: Python, SQL, TypeScript, React, Next.js, FastAPI, Docker, Kubernetes, Git

Cloud & Data: Azure, Pandas, NumPy, Scikit-learn, PostgreSQL, MongoDB, SQL Server, Power BI